

# Game Changer: Software-Defined Storage and Container Schedulers

David vonThenen  
{code} by Dell EMC  
@dvonthenen  
dvonthenen.com  
github.com/dvonthenen



# Agenda

- Container Schedulers
- Containers In Production
- Software-Defined Storage (SDS)
- Schedulers + SDS = Game Changing
- Demo



# Schedulers



# What is a Scheduler?

- Fair and efficient workload placement
- Adhering to a set of constraints
- Quickly (and deterministically) dispatching jobs
- Robust and tolerates errors




# Let's take a look: Apache Mesos

- Is a Container Scheduler
  - Docker
  - Unified Containerizer
- Cluster Manager
- Task placement based on CPU, Memory, and Disk
- User defined constraints
- 2 Layer Scheduler – Offer/Accept Model

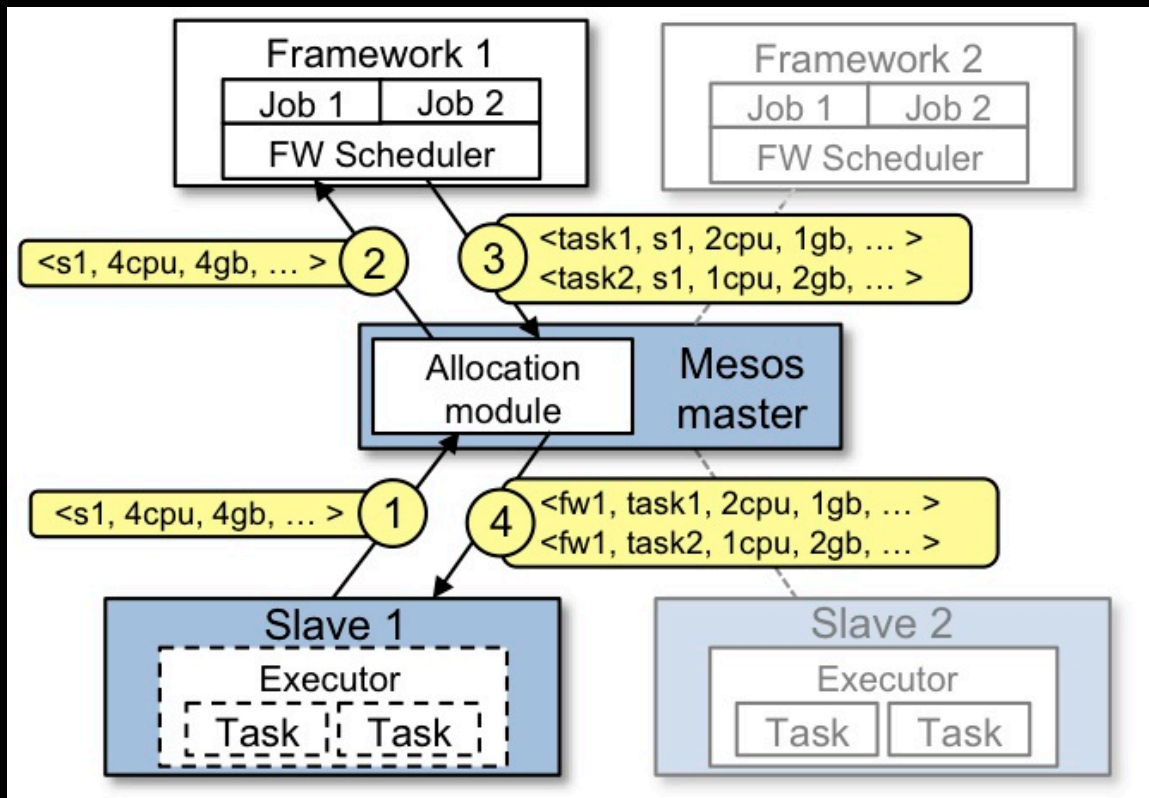


# Mesos Frameworks

- Ability to sub-schedule tasks based on Application needs
- Framework implements a Scheduler and Executor
  - Scheduler – Accepts/Denies resources
  - Executor – Application
- Multiple Frameworks run within the cluster



# Framework / Offer Mechanism















# Containers






# Containers Today

- Many container workloads are long running
- Many have state: user data, configuration, and etc
- Top 7 of 12 Apps in Docker Hub are persistent applications

 nginx official	3.1K STARS	10M+ PULLS	> DETAILS
 busybox official	672 STARS	10M+ PULLS	> DETAILS
 ubuntu official	4.0K STARS	10M+ PULLS	> DETAILS
 registry official	845 STARS	10M+ PULLS	> DETAILS
 swarm official	346 STARS	10M+ PULLS	> DETAILS
 redis official	2.2K STARS	10M+ PULLS	> DETAILS
 mongo official	1.9K STARS	10M+ PULLS	> DETAILS
 mysql official	2.4K STARS	10M+ PULLS	> DETAILS
 node official	2.2K STARS	10M+ PULLS	> DETAILS
 postgres official	2.1K STARS	10M+ PULLS	> DETAILS
 elasticsearch official	1.2K STARS	10M+ PULLS	> DETAILS
 wordpress official	1.0K STARS	5M+ PULLS	> DETAILS



# Death of a Container

- Where does my data go?
- Turned to the compute node's local disk to store data
- What happens on a node failure?
- Production applications require high availability
- External Storage!



# How do we achieve this?

- REX-Ray
  - Vendor agnostic storage orchestration engine
  - AWS, GCE, ScaleIO, VirtualBox, many more
  - GitHub: <https://github.com/emccode/rexray>
- mesos-module-dvdi
  - Provides hooks to Mesos agent nodes to manage external storage
  - GitHub: <https://github.com/emccode/mesos-module-dvdi>



# Enablement is Out-Of-Band

- The “glue” that combines compute to external storage is add-on to the resource manager
- Obvious but easily dismissive answer: DevOps
  - Software upgrades? On all nodes...
  - Maintenance? Infrastructure, Storage Platform, etc
  - Changes to Container Scheduler? Behaviors, APIs, etc
- Just make it happen!
- Almost 100% of the way there...



Got to be an easier way...



# Software-Defined Storage



# What are they?

- Software-Defined Storage (SDS) serve as abstraction layer above underlying storage
- Provides a (programmatic) mechanism to provision storage
- Varying degrees of SDS: NFS, VMware Virtual Volumes



# What makes them unique?

- Manage provisioning and data independent of underlying hardware (operational)
- Abstract consumed logical storage from underlying physical storage (physical)
- Automation of policy driven SLAs both external (users) and internal (platform)





# Let's take a look: ScaleIO

- Scale-out block storage
- Linear performance
- Elastic architecture
- Infrastructure agnostic
- Try ScaleIO as a free download:


<https://www.emc.com/products-solutions/trial-software-download/scaleio.htm>

ScaleIO



# Scale-out Block Storage

- Scale from 3 nodes to 1000s of nodes
- Add storage services and servers on the fly to increase capacity and performance
- Storage growth always automatically aligned with application needs



# Elastic Architecture



Add, remove, re-allocate, on the fly, without stopping IO

AUTO-REBALANCE

when resources are added

AUTO-REBUILD

when resources fail or removed

**NO CAPACITY PLANNING OR MIGRATION!**



# Infrastructure Agnostic

Mix and match OS, hypervisors, platforms, media in the same ScaleIO system



OPERATING  
SYSTEM



HYPERVERSORS



CLOUD


HDDs



SSDs



PCIe  
Flash



MEDIA




# Game Changer



# Let's Review...

- Container Schedulers:
  - Great platform for container management
  - Needs persistent storage for production Apps
  - Adding persistent storage out-of-band presents challenges
- Software-Defined Storage:
  - Scale-out storage
  - Elastic architecture
  - Infrastructure agnostic



Schedulers + SDS = ????

One more thing...



# Better than the Sum of Our Parts

- Let's create a Software-Defined Storage Framework
- ScaleIO + Mesos Framework = Awesome Sauce!
- <https://github.com/codedellemc/scaleio-framework>





# SDS Framework = Mind Blown

- External persistent storage **native** to scheduling platform
- Globally accessible storage
- Storage array? Reduce complexity
- Reduces maintenance
- Deploy Anywhere!



# What this Means for your Apps

- No data loss on infrastructure failure
- Insulates changes with cluster manager (APIs, etc)
- Highly Available containers and Apps!
- Production ready!
- Tolerates failures



# Surprising Combination



# Demo



# Configuration

- 3 Node Mesos Cluster (Management)
- 2 Mesos Agent nodes (Compute)
  - Initially the first node online
  - Second node will be onboarded or introduced later
- ScaleIO Cluster (Scale-out storage)
  - 3 management nodes
  - 180 GB local disks on each management node to comprise this storage pool




# Configuration (Cont.)

- ScaleIO Framework
  - GitHub: <https://github.com/codedellemc/scaleio-framework>
- Persistent External Storage
  - Using REX-Ray
    - › GitHub: <https://github.com/emccode/rexray>
  - Using mesos-module-dvdi
    - › GitHub: <https://github.com/emccode/mesos-module-dvdi>



# The Moving Parts



#CodeOpen



# Demo

[codedellemc.com](http://codedellemc.com)




#CodeOpen



# Thank you

[codedellemc.com](http://codedellemc.com)

#CodeOpen



[codedellemc.com](http://codedellemc.com)

**D~~E~~LL EMC**