

# How Container Schedulers and Software-Defined Storage will Change the Cloud

David vonThenen  
{code} by Dell EMC  
@dvonthenen  
<http://dvonthenen.com>  
[github.com/dvonthenen](https://github.com/dvonthenen)



# Agenda

- Review of Software-Defined Storage
- Container Schedulers
- Schedulers + Software-Defined Storage = Awesome!
- To the Cloud!!
- Demo



# Review of Software-Defined Storage



# What are they?

- Many definitions... most agree on:
- Software-Defined Storage (SDS) serve as abstraction layer above underlying storage
  - Installed software versus Physical Storage Arrays
- Provides a (programmatic) mechanism to provision storage
- Varying degrees of SDS: NFS, VMware VSAN

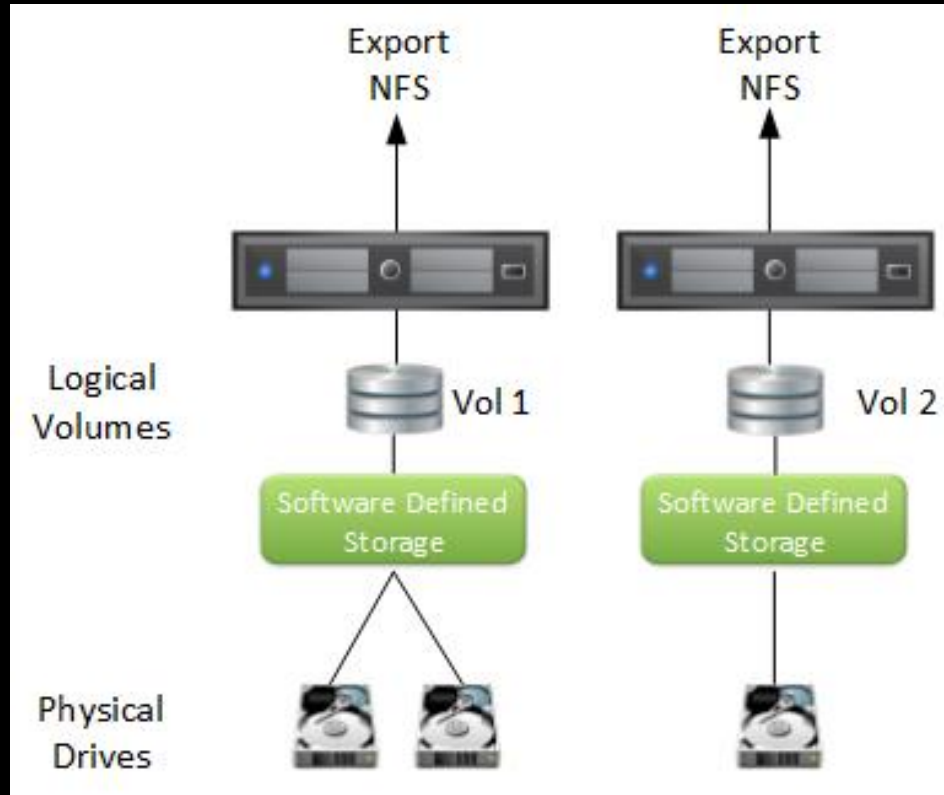


# What makes them unique?

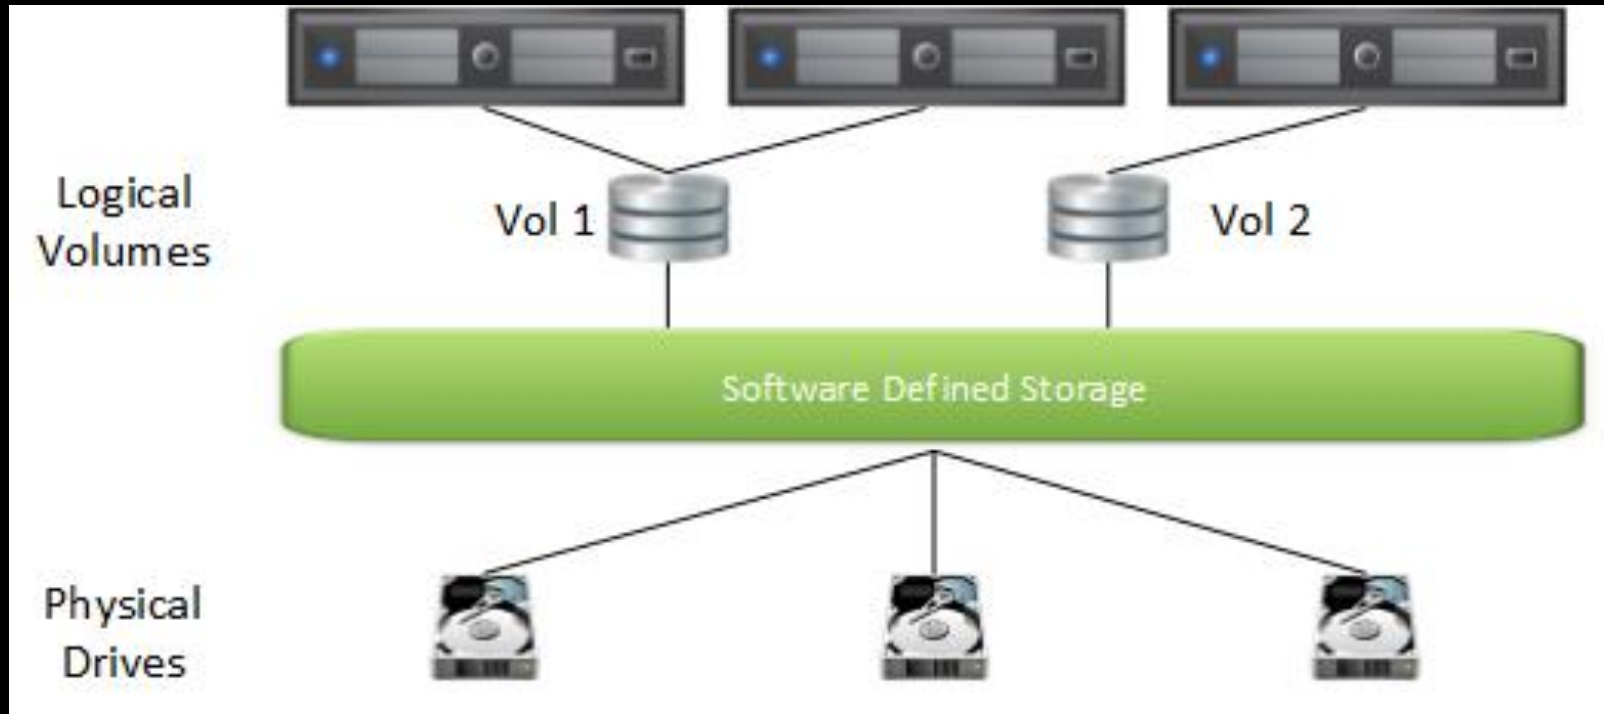
- Operational - Manage provisioning and data independent of underlying hardware
- Physical - Abstract consumed logical storage from underlying physical storage
- Policy - Automation of policy driven both external (users) and internal (platform)
- Day 2 Operations - Maintenance is inherently different



# What that might look like...



# What that might look like...



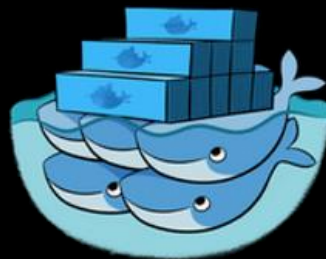
# Container Schedulers






# What is a Scheduler?

- Fair and efficient workload placement
- Adhering to a set of constraints
- Quickly (and deterministically) dispatching jobs
- Robust and tolerates errors



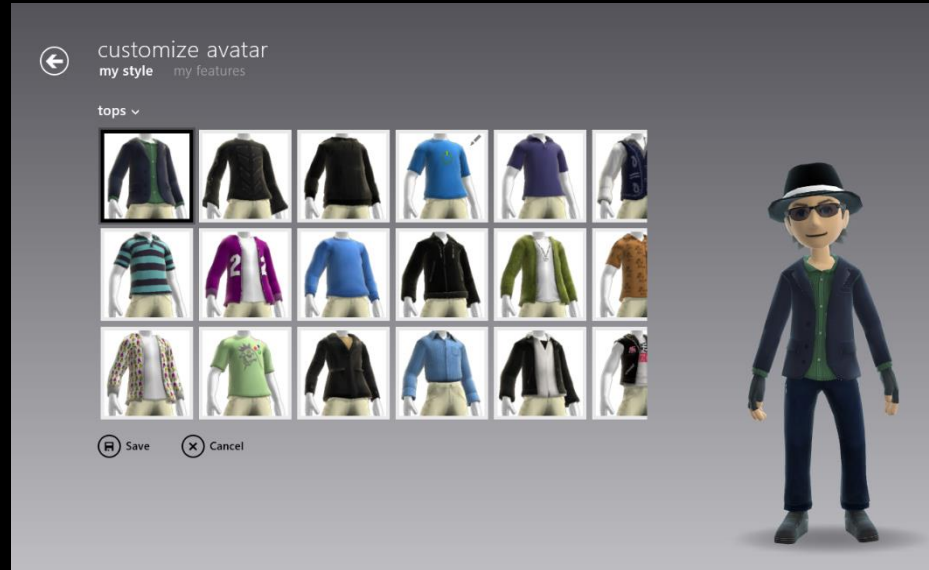
# Scheduling Work

- Containers like...
  - Docker
  - Mesos Unified Containerizer
  - rkt (CoreOS)
- Cluster Manager
- Task placement based on resource
- Operational constraints



# Custom Scheduling

- Many allow creation of own custom Scheduler
- Customization for your application:
  - Run-Time?
  - Availability?
  - Fault Tolerance?
  - Hardware Accel?
  - Location?



# Apache Mesos

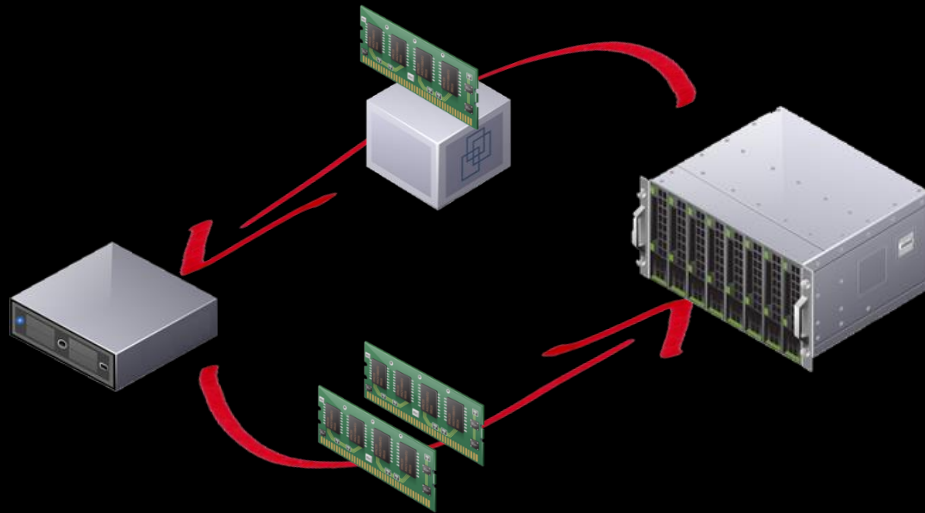


More

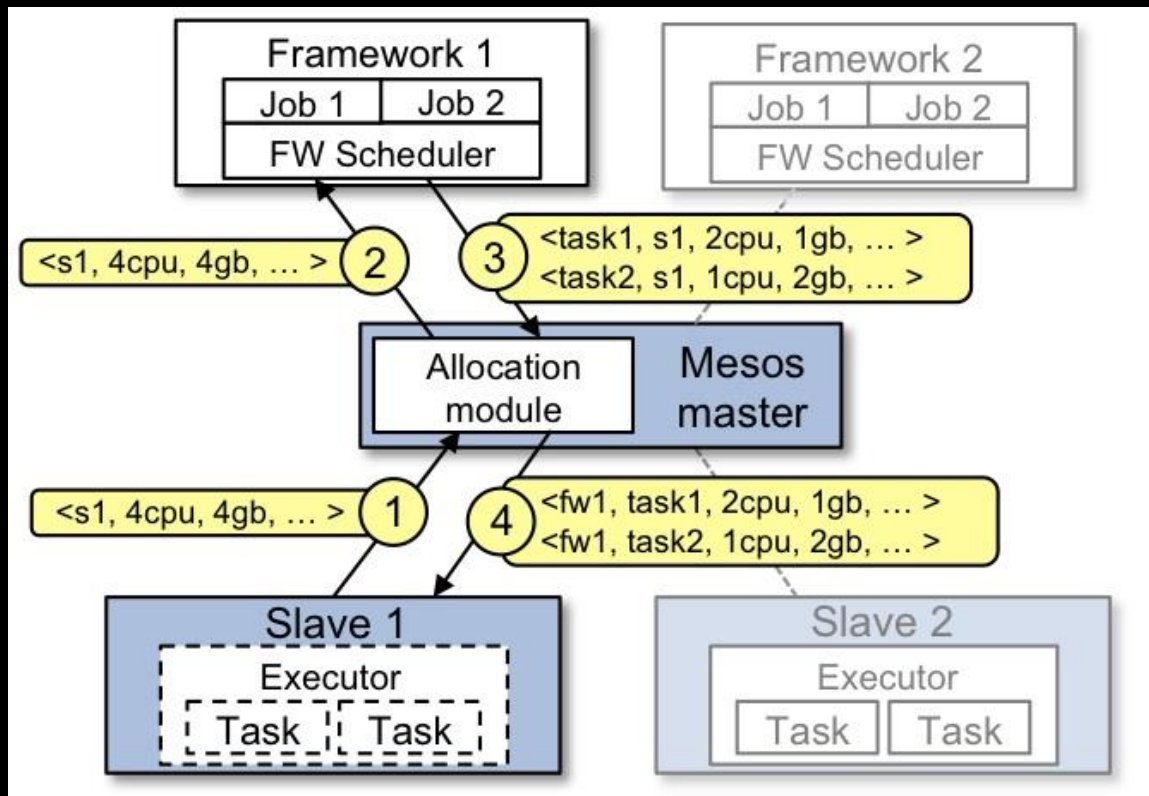


# Mesos Frameworks

- Ability to schedule tasks based on Application needs
- Framework implements a Scheduler and Executor
  - Scheduler – Accepts/Denies resources
  - Executor – Application
- Offer / Accept Mechanism
- Multiple Frameworks run within the cluster



# Framework / Offer Mechanism



# Schedulers and Software Defined Storage



# Better Together

- Let's create a Software-Defined Storage Framework
- ScaleIO + Mesos Framework = Awesome Sauce!
- First released in Sept 2016.  
Now on version 0.3.1
- <https://github.com/codedellemc/scaleio-framework>





# Let's take a look: ScaleIO

- Scale-out block storage
- Linear performance
- Elastic architecture
- Infrastructure agnostic
- Try ScaleIO. It's a free download!

ScaleIO

<https://www.emc.com/products-solutions/trial-software-download/scaleio.htm>















# SDS Framework = Mind Blown

- Framework installs and configures Storage Platform on all Scheduler's compute nodes
- Persistent storage **native** to scheduling platform
- Globally accessible storage
- What Storage array? Reduce complexity
- Deploy Anywhere!



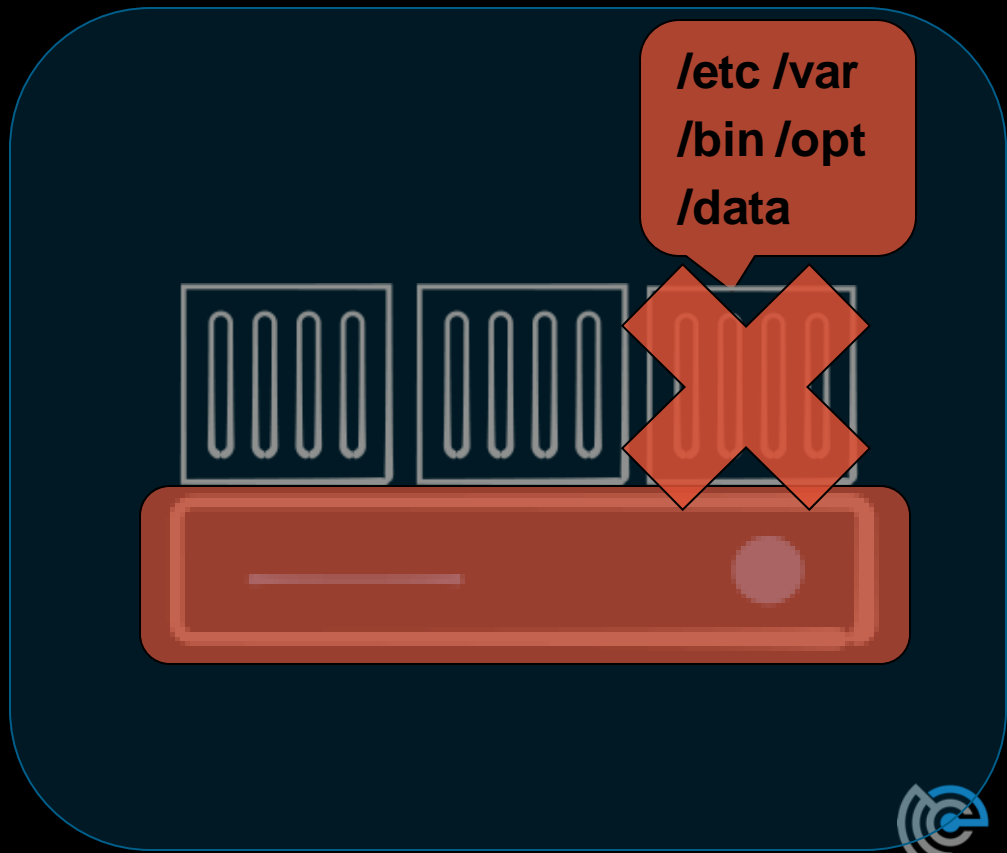
# Containers Today

- Many container workloads are long running
- Many have state: user data, configuration, and etc
- Top 7 of 12 Apps in Docker Hub are persistent applications

 <b>nginx</b> official	3.1K STARS	10M+ PULLS	<a href="#">➤</a> DETAILS
 <b>busybox</b> official	672 STARS	10M+ PULLS	<a href="#">➤</a> DETAILS
 <b>ubuntu</b> official	4.0K STARS	10M+ PULLS	<a href="#">➤</a> DETAILS
 <b>registry</b> official	845 STARS	10M+ PULLS	<a href="#">➤</a> DETAILS
 <b>swarm</b> official	346 STARS	10M+ PULLS	<a href="#">➤</a> DETAILS
 <b>redis</b> official	2.2K STARS	10M+ PULLS	<a href="#">➤</a> DETAILS
 <b>mongo</b> official	1.9K STARS	10M+ PULLS	<a href="#">➤</a> DETAILS
 <b>mysql</b> official	2.4K STARS	10M+ PULLS	<a href="#">➤</a> DETAILS
 <b>node</b> official	2.2K STARS	10M+ PULLS	<a href="#">➤</a> DETAILS
 <b>postgres</b> official	2.1K STARS	10M+ PULLS	<a href="#">➤</a> DETAILS
 <b>elasticsearch</b> official	1.2K STARS	10M+ PULLS	<a href="#">➤</a> DETAILS
 <b>wordpress</b> official	1.0K STARS	5M+ PULLS	<a href="#">➤</a> DETAILS

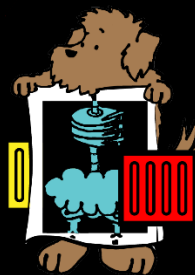
# Death of a Container

- Where does my data go?
- Turned to the compute node's local disk to store data
- What happens on a node failure?
- Production applications require high availability
- External Storage!



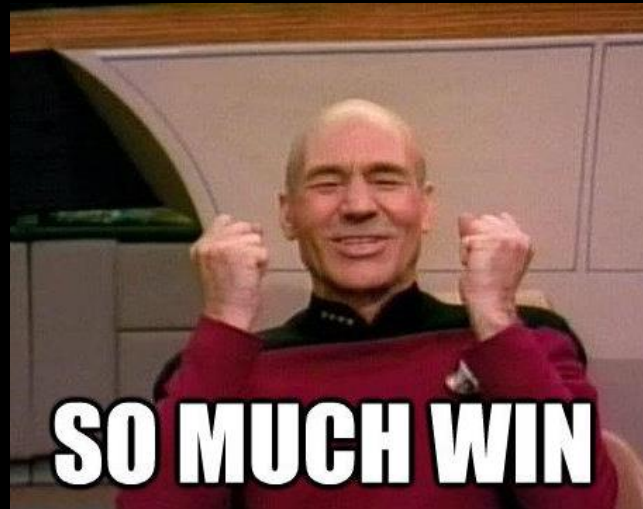
# External Storage Enablement

- REX-Ray
  - Vendor agnostic storage orchestration engine
  - AWS, GCE, ScaleIO, VirtualBox, many more
  - <https://github.com/codedellemc/rexray>
- mesos-module-dvdi
  - Hook for Mesos nodes to manage external storage
  - <https://github.com/codedellemc/mesos-module-dvdi>
  - Contributed back to and is apart of Mesos proper



# What this Means for your Apps

- Tolerates node failures
- Highly Available containers and Apps!
- Insulates changes with:
  - container scheduler (APIs, etc)
  - storage platform (workflows, APIs, etc)
- Production ready!



# To the Cloud!



# Moving towards the Cloud

- Applications with management APIs
- Cloud is perfect to enable DevOps
- What makes these cloud accessible?





# Self Monitoring Apps

- Framework deploy and configure applications.
- Enable application monitoring via Management APIs
- Determine health and remediate!
- Can fix themselves, but to what end?



# Self-aware Applications

- AWS SDK – 10 Language bindings
- Software-based Storage Platform with a Cloud Platform driven by APIs
- Applications that change their environment
  - Maintenance, Remediation, Performance, etc
- Self-aware applications! Skynet!



# Premise: Self Managing

- Framework can monitor and self remediate Software-based Storage Platform
- The Scenario:
  - ScaleIO has a Storage Pool that is approaching full
  - Identifies the health check warning
  - Creates new EBS volumes in EC2 to expand the Storage Pool



# Demo



# Configuration

- Mesos Configuration
  - 3 Node Mesos Cluster (Management)
  - 3 Mesos Agent nodes (Compute)
- ScaleIO Cluster (Scale-out storage)
  - Will install on top of 3 Mesos Agent nodes
  - 180 GB local disks on each node to make up this Storage Pool

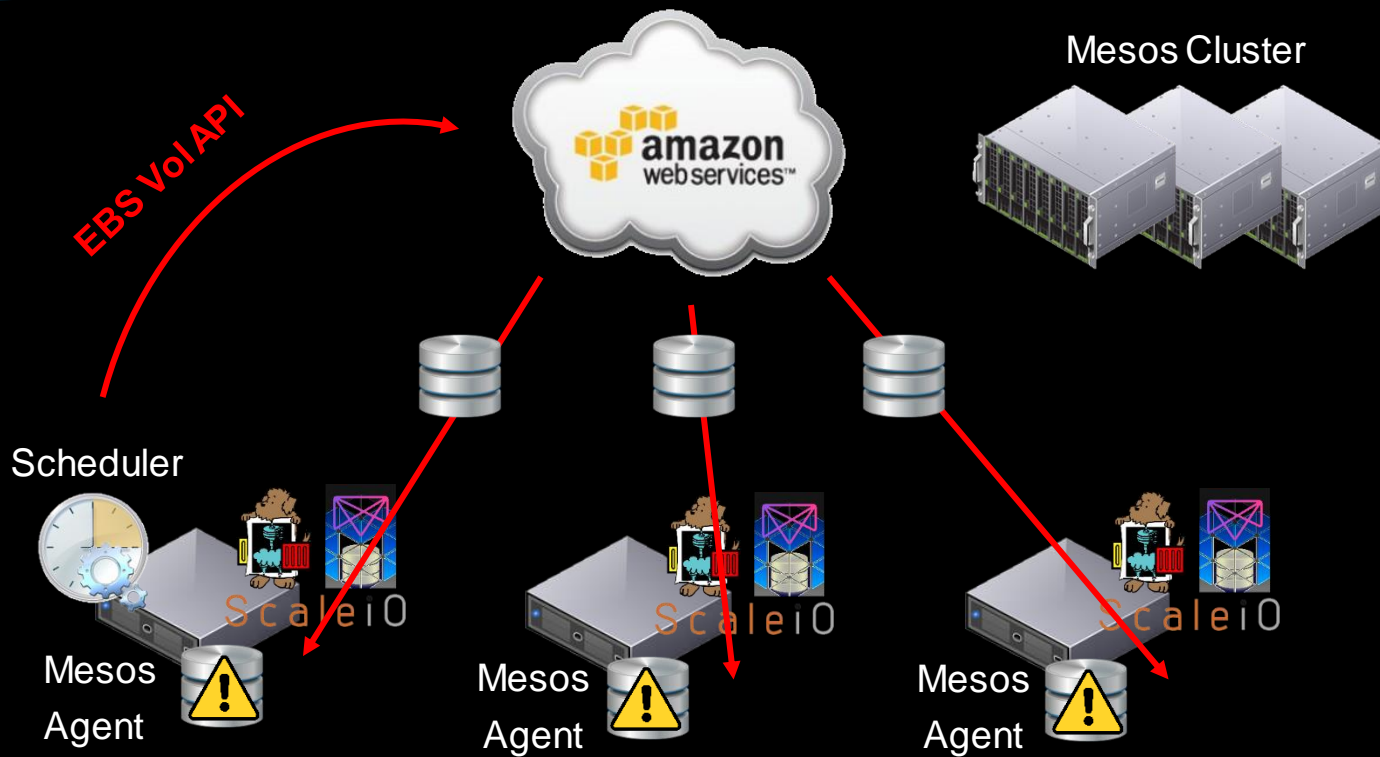


# Configuration (Cont.)

- ScaleIO Framework
  - GitHub: <https://github.com/codedellemc/scaleio-framework>
- Persistent External Storage
  - Using REX-Ray
    - › GitHub: <https://github.com/emccode/rexray>
  - Using mesos-module-dvdi
    - › GitHub: <https://github.com/emccode/mesos-module-dvdi>



# The Moving Parts



#CodeOpen

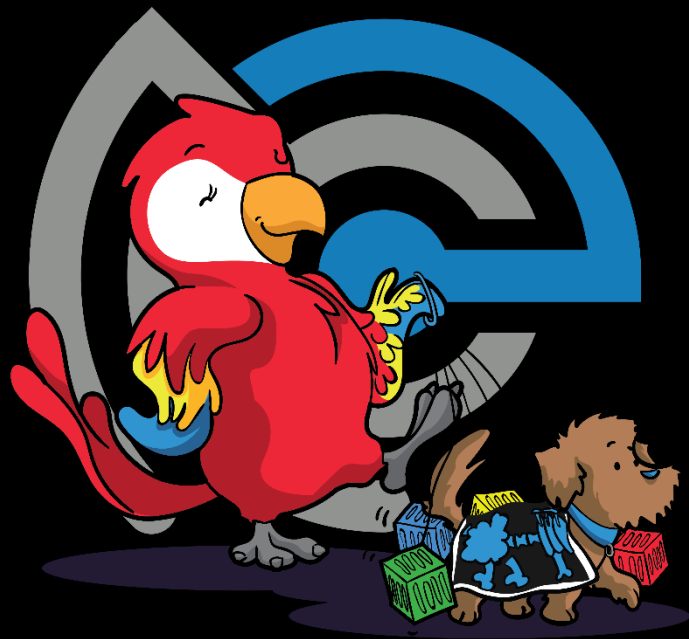


# Thank you

[codedellemc.com](http://codedellemc.com)



#CodeOpen



[codedellemc.com](http://codedellemc.com)